

Magellan: Spark as a Geospatial Analytics Engine

Ram Sriharsha (@halfabrane)

Product Manager, Apache Spark @ Databricks





Who Am I?

- Product Manager of Apache Spark @ Databricks
- PMC Member, Committer, Apache Spark
- Prior to Databricks
 - Hortonworks Architect, Spark and Data Science
 - Magellan, Geospatial Analytics on Spark
 - Yahoo Labs, Principal Research Scientist in Scalable Machine Learning
 - Login Risk Detection, Search Advertising Click Prediction, Online Clustering/ Classification.

Agenda

- What is Geospatial Analytics?
- The basic operations in Magellan
- Some geometric algorithms used by Magellan
- Internals: How Magellan works with Spark SQL
- Upcoming work: Spatial Indices

What is geospatial analytics?



How do pickup/ dropoff neighborhood hotspots evolve with time?

Correct GPS errors with more
Accurate landmark measurements


Incorporate location in IR and search
advertising

Do we need one more library?


- Spatial Analytics at scale is challenging
 - Single machine libraries not fast enough
 - No scalable implementations exist
- Ancient Data Formats
 - Do not leverage columnar storage, metadata hard to parse and index
 - No spatial indexing
- Geospatial Analytics is not simply BI anymore
 - (approx) Near neighbor queries
 - Map matching

The basic operations

Introduction to Magellan



Introduction to Magellan




```
Polygon = (  
  [0, 5],  
  [(0.0, 0.0),(1.0, 0.0),  
   (1.0, 2.0),(0.0, 2.0),  
   (0.0, 0.0),  
   (0.3, 0.3),  
   (0.6, 0.3),  
   (0.6, 0.9),  
   (0.3, 0.9),  
   (0.3, 0.3)  
])
```


Reading from common data formats

polygon	metadata
([0], [(-122.4413024, 7.8066277), ...])	neighborhood -> Marina
([0], [(-122.4111659, 37.8003388), ...])	neighborhood -> North Beach

```
sqlContext.read.format("magellan")  
.load(${neighborhoods.path})
```



```
sqlContext.read.format("magellan")  
.option("type", "geojson")  
.load(${neighborhoods.path})
```



Geometric Expressions

polygon	metadata
[[0], [[-122.4413024, 7.8066277], ...]]	neighborhood -> Marina
[[0], [[-122.4111659, 37.8003388], ...]]	neighborhood -> North Beach



polygon	metadata
[[0], [[-122.4111659, 37.8003388], ...]]	neighborhood -> North Beach

```
neighborhoods.filter(
```

```
  point(-122.4111659, 37.8003388)
```

```
  within
```

```
    'polygon
```

```
  ).show()
```

Shape literal

Boolean Expression

Spatial Joins

polygon	metadata
([0], [(-122.4111659, 37.8003388), ...])	neighborhood -> North Beach
([0], [(-122.4413024, 7.8066277), ...])	neighborhood -> Marina

point
(-122.4111659, 37.8003388)
(-122.4343576, 37.8068007)

`points.join(neighborhoods).
where('point within 'polygon').
show()`

point	polygon	metadata
(-122.4343576, 37.8068007)	([0], [(-122.4111659, 37.8003388), ...])	neighborhood -> North Beach

Near neighbor queries

polygon	metadata
([0], [(-122.4111659, 37.8003388), ...])	neighborhood -> North Beach
([0], [(-122.4413024, 7.8066277), ...])	neighborhood -> Marina



point	polygon	metadata
(-122.4343576, 37.8068007)	([0], [(-122.4111659, 37.8003388), ...])	neighborhood -> North Beach

```
neighborhoods.filter(  
    point(-122.4111659, 37.8003388).buffer(0.1)  
    intersects  
    'polygon'  
)show()
```

Advantage of embedding geometric queries in Spark SQL


```
val uber = sqlContext.read.format("csv").path(${uber.path})

val neighborhoods = sqlContext.read.format("magellan").
    load(${neighborhoods.path}).
    select($"polygon", $"metadata"). cache()

Val joined = neighborhoods.
    join(uber).
    where($"point".within($"polygon")).
    select($"tripId", $"timestamp", explode($"metadata").as(Seq("k", "v"))).
    withColumnRenamed("v", "neighborhood")
```

algorithms

`point within` polygon



The Join

Inherits all join optimizations from Spark SQL

- if neighborhoods table is small, Broadcast Cartesian Join
- else Cartesian Join

status

- Magellan 1.0.3 available as Spark Package.
- Scala , Spark 1.4
- Github: <https://github.com/harsha2010/magellan>
- Blog: <http://hortonworks.com/blog/magellan-geospatial-analytics-in-spark/>
- Notebook example: <http://bit.ly/1GwLyrV>
- Input Formats: ESRI Shapefile, GeoJSON, OSM-XML
- Please try it out and give feedback!

- Magellan 1.0.4 release upcoming in end of June.
- Preview available in Databricks end of May.
- Spark 1.6, 2.0
- Python, Scala
- Tight integration with Tungsten' s memory layout
- Codegen for all operators
- Supports within, contains, intersects, shape literals, near neighbor queries

The internals

Shapes as Data Types

- Points, Polygons, Lines, Polylines are Spark SQL Data Types (**UserDefinedType**)
- Tungsten Encoding:
 - 8 bit type indicator (1 = point, 3 = polyline, 5 = polygon, ...)
 - $16 * 4$ bit bounding box (xmin, ymin, xmax, ymax)
 - For point, x and y coordinates = $2 * 16$ bits
 - For polygon, indices, coordinates arrays = $8 * \# \text{ of rings} + 16 * \# \text{ of points} * 2$
 - ...

Data Sources

- `SpatialRelation` extends `BaseRelation`, `PrunedFilteredScan`
- `GeoJSONRelation`, `ShapeFileRelation`, `OSMRelation` for GeoJSON, ESRI ShapeFile, OSMXML
- Pushes predicates and filters down if possible
- Returns `Row[point, polygon, polyline, meta]`
- `Metadata = Map[String, String]`

`point within` polygon

- Create a case class **Within** that extends **BinaryExpression**
- Override **genCode(ctx: CodeGenContext, ev: GeneratedExpressionCode)** to return generated code
 - Generated code optimizes by bounding box intersections/contains
 - Takes advantage of Tungsten format
- Make use of ctx to store expensive initialization and reusable objects
- Use implicit conversions to and from **Column/**

Within codegen

```
nullSafeCodeGen(ctx, ev, (c1, c2) => {  
    s"" +  
    s"Double lxmin = $c1.getDouble(1);" +  
    s"Double lymin = $c1.getDouble(2);" +  
    s"Double lxmax = $c1.getDouble(3);" +  
    s"Double lymin = $c1.getDouble(4);" +  
    s"Double rxmin = $c2.getDouble(1);" +  
    s"Double rymin = $c2.getDouble(2);" +  
    s"Double rxmax = $c2.getDouble(3);" +  
    s"Double rymax = $c2.getDouble(4);" +  
    s"Boolean within = false;" +  
    s"if (rxmin <= lxmin && rymin <= lymin && rxmax >= lxmax && rymax >= lymin) {" +  
    s"Integer ltype = $c1.getInt(0);" +  
    s"Integer rtype = $c2.getInt(0);" +  
    s"magellan.Shape leftShape = (magellan.Shape)" +  
    s"    ((org.apache.spark.sql.types.UserDefinedType<magellan.Shape>)" +  
    s"    serializers.get(ltype)).deserialize($c1);" +  
    s"magellan.Shape rightShape = (magellan.Shape)" +  
    s"    ((org.apache.spark.sql.types.UserDefinedType<magellan.Shape>)" +  
    s"    serializers.get(rtype)).deserialize($c2);" +  
    s"within = rightShape.contains(leftShape);" +  
    s"}" +  
    s"${ev.value} = within;" +  
    s"}  
})
```


The next steps



The join revisited

What is the time complexity?

- m points, n polygons (assume average k edges)
- l partitions
- $O(mn/l)$ computations of 'point within 'polygon
- $O(ml)$ communication cost
- Each 'point within 'polygon costs $O(k)$
- Total cost = $O(ml) + O(mnk/l)$

=> $O(m\sqrt{n}\sqrt{k})$ cost, with $O(\sqrt{n}\sqrt{k})$ partitions

Optimization

Do we need to send every point to every partition?

Do we need to compute 'point in 'neighborhood for each neighborhood within a given partition?


2D Indices

- Quad Trees
- R Trees
- Dimensional Reduction
 - Hashing
 - PCA
 - Space Filling Curves


What does a good dimensional reduction need?

- Preserve (approximate) nearness in ambient space
- Enable range queries
- Little/ no collision


Row Order Curve




Snake Order Curve




Z Order Curve



Binary Representation



Binary Representation



Properties

- Locality: two points differing in small # of bits => nearness
 - Converse not necessarily true
- Containment
- Efficient construction
- Nice bounds on precision

If you are familiar with GeoHash

Its nothing but a Z Order Curve!

*Start with Bounding Box = $(-180, -90, 180, 90)$ and
compute Binary Representation.*

Then convert to Base 32 encoded String

You Obtain the GeoHash

How to speed up join?

- Preprocess points:
 - Index each point to a unique geohash
- Preprocess polygons:
 - Index each polygon to a set of geohashes
- Inner join on geohash
- Filter out edge cases

Spark Implementation


- Define **SpatialJoinStrategy** that extends **org.apache.spark.sql.Strategy**
 - Add logic to decide when to trigger this join
 - Only trigger if geospatial queries
 - Only trigger if join is complex: if $n \sim O(1)$ then broadcast join is good enough
- Override **BinaryNode** to handle the physical execution plan ourselves
 - Override `execute(): RDD` to execute join and return results
- Stitch it up using **ExperimentalStrategies** in **SQLContext**

Persistent Indices

Often the geometries do not change (or change slowly)

Can we pre index them?

Overall architecture



Contributions to Magellan welcome!

- Algorithms
 - Map Matching
 - Persistent Indices
- Integration with Spark
 - Python API, R API?
 - Encoders
- Data Formats