

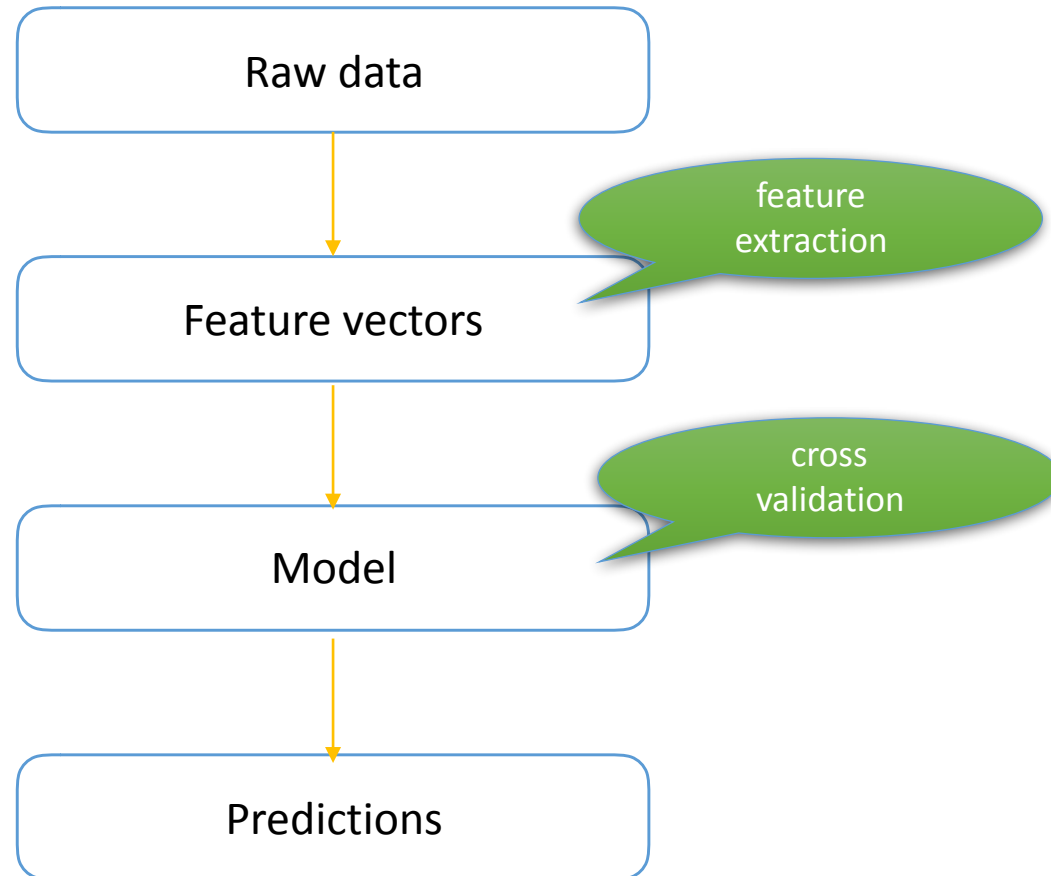
# Beyond MLlib: Scale up Advanced Machine Learning on Spark

Yizhi Liu

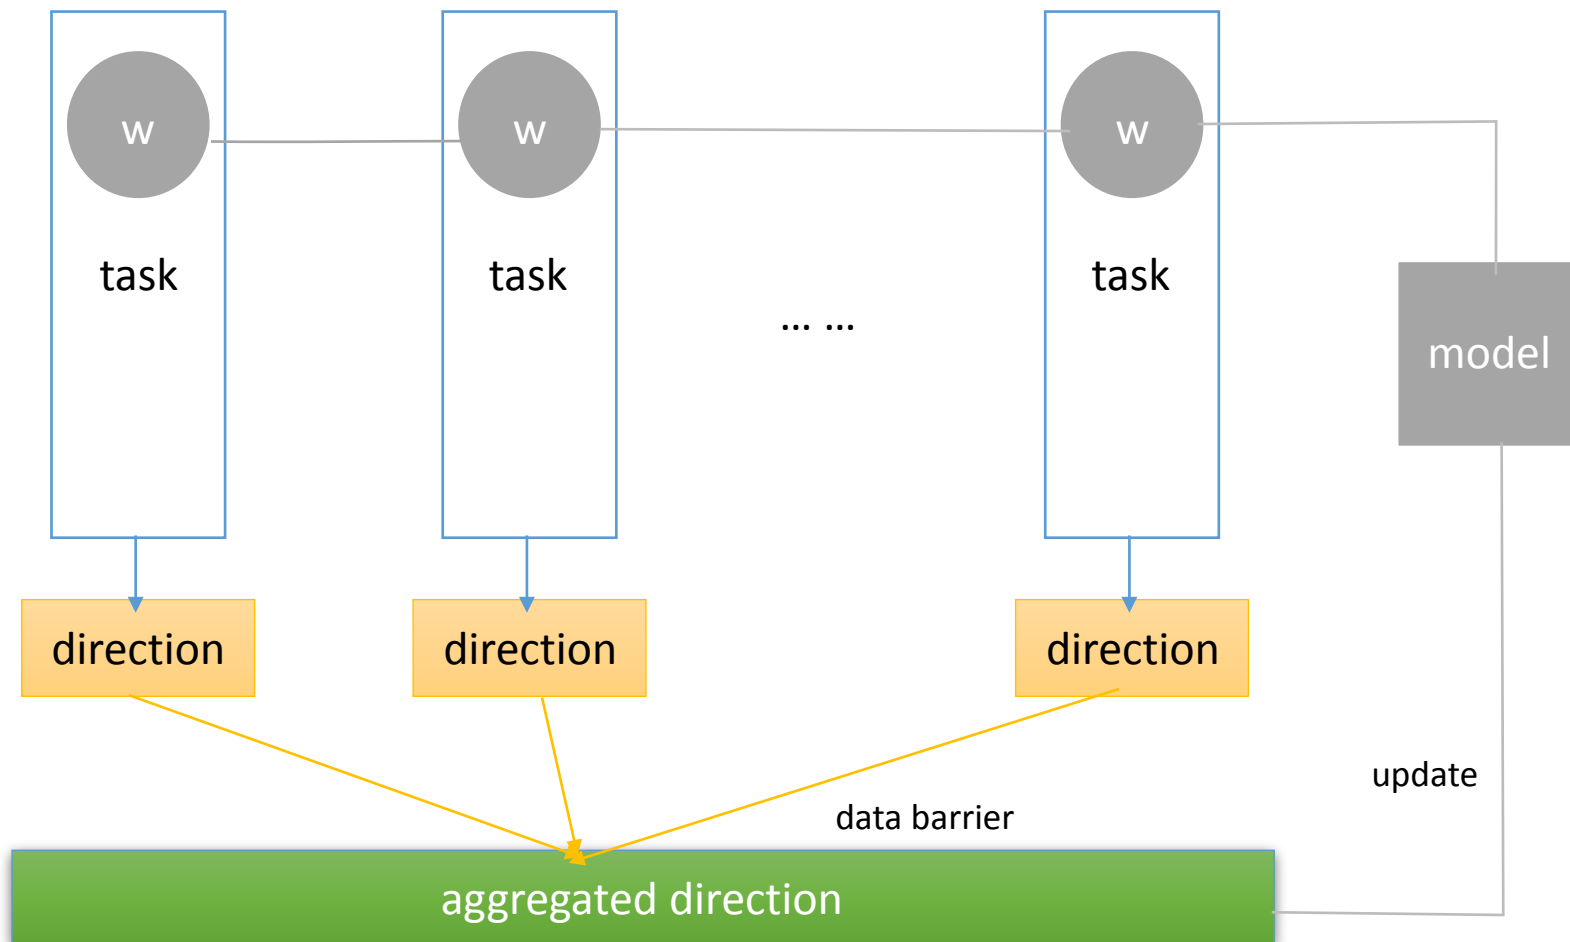
Qihoo 360 / MVAD  
DMLC member  
MXNet committer

# Algorithms in Spark ML

- Classification and Regression
  - Linear models
  - Naive Bayes
  - Decision Trees
  - Ensembles of Trees
  - Isotonic Regression
- Collaborative Filtering
  - ALS
  - ... ..
- Clustering
  - Latent Dirichlet Allocation
  - ... ..
- Dimensionality Reduction
  - ... ..



# Computational Model



## Advanced ML Algorithms


dmlc  
**XGBoost**

XGBoost is an optimized distributed gradient boosting library designed to be highly **efficient**, **flexible** and **portable**


MXNet is a deep learning framework designed for both **efficiency** and **flexibility**

dmlc  
**mxnet**


# Deep Learning



"deep learning" trend in the past 10 years

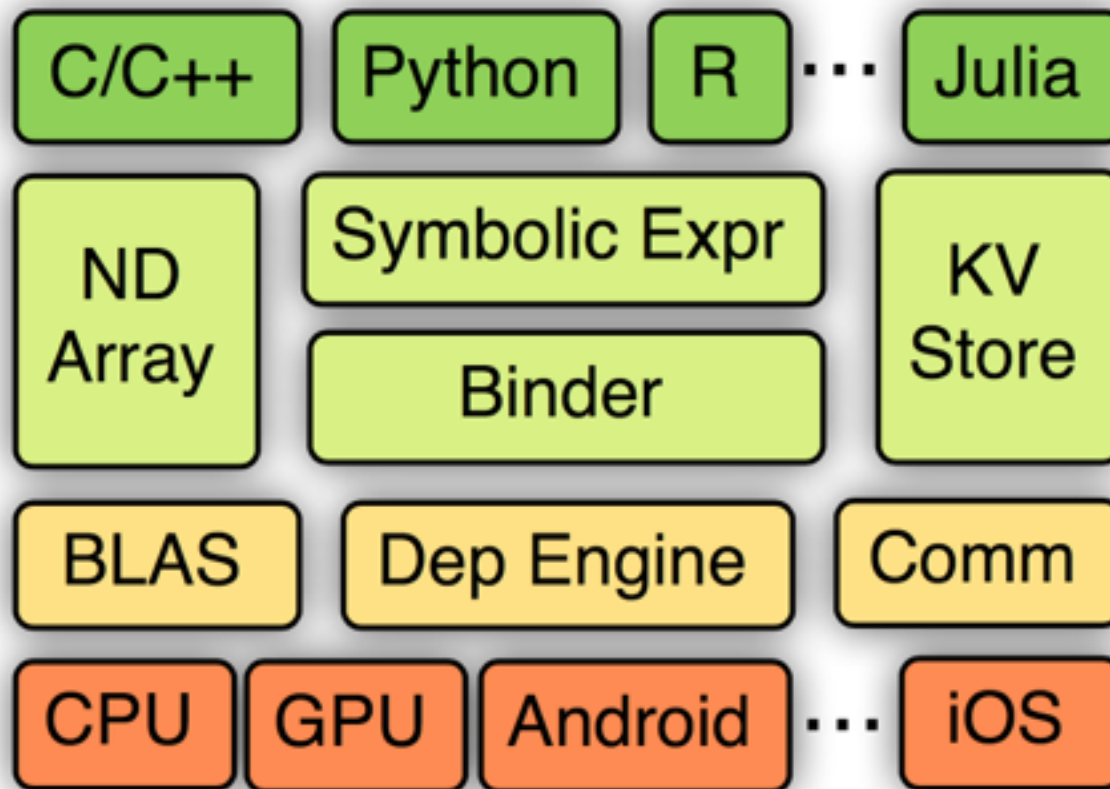


# Deep Learning



- Hard to define the network
- Huge computational cost
  - Convolution layers
  - Fully connected layers
- Memory limit
- Way to distribute training process

## MXNet Overview





# Train Deep Network on MXNet

```
// model definition
```

```
val data = Symbol.Variable("data")
val fc1 = Symbol.FullyConnected(name = "fc1")(Map("data" -> data, "num_hidden" -> 128))
val act1 = Symbol.Activation(name = "relu1")(Map("data" -> fc1, "act_type" -> "relu"))
val fc2 = Symbol.FullyConnected(name = "fc2")(Map("data" -> act1, "num_hidden" -> 64))
val act2 = Symbol.Activation(name = "relu2")(Map("data" -> fc2, "act_type" -> "relu"))
val fc3 = Symbol.FullyConnected(name = "fc3")(Map("data" -> act2, "num_hidden" -> 10))
val mlp = Symbol.SoftmaxOutput(name = "sm")(Map("data" -> fc3))
```

(declaratively)  
define layers

```
// setup model and fit the training set
```

```
val model = FeedForward.newBuilder(mlp)
    .setContext(Context.cpu())
    .setNumEpoch(10)
    .setOptimizer(new SGD(learningRate = 0.1f, momentum = 0.9f, wd = 0.0001f))
    .setTrainData(trainDataIter)
    .setEvalData(valDataIter)
    .build()
```

set devices here, e.g.,  
Context.gpu(0,1,2,3)

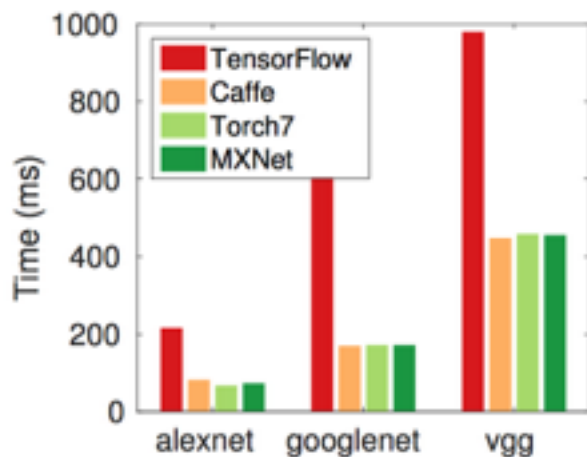
```
val probArrays = model.predict(valDataIter)
// in this case, we do not have multiple outputs
require(probArrays.length == 1)
val prob = probArrays(0)
// get predicted labels
val py = NDArray.argmaxChannel(prob)
// deal with predicted labels py
```

user defined optimizer





# MXNet Benchmarks

Compare MXNet to others on a single forward-backward performance



Internal memory usage of MXNet under various allocation strategies



# GPU Support

NDArray is imperative

specify device here

```
val weight = NDArray.empty(Shape(3, 2), Context.gpu(0))  
weight -= eta * (grad + lambda * weight);
```

```
val weightOnCpu = weight.copyTo(Context.cpu())
```

Copy to another device

One code for both CPU and GPU,  
(mshadow) translates at compile time

```
template<typename xpu>  
void UpdateSGD(Tensor<xpu, 2> weight, const Tensor<xpu, 2> &grad,  
               float eta, float lambda) {  
    weight -= eta * (grad + lambda * weight);  
}
```

## GPU Support: mshadow

<https://github.com/dmlc/mshadow>

- Efficient: all the expression you write will be lazily evaluated and compiled into optimized code.
  - No temporal memory allocation will happen for expression you write.
  - mshadow will generate specific kernel for every expression you write in compile time.
- Device invariant: you can write one code and it will run on both CPU and GPU.
- Simple: mshadow allows you to write machine learning code using expressions.
- Whitebox: put a float\* into the Tensor struct and take the benefit of the package, no memory allocation is happened unless explicitly called.
- Lightweight library: light amount of code to support frequently used functions in machine learning.
- Extendable: user can write simple functions that plugs into mshadow and run on GPU/CPU, no experience in CUDA is required.
- MultiGPU and Distributed ML: mshadow-ps interface allows user to write efficient MultiGPU and distributed programs in an unified way.

# Distributed Training

```
val envs = Map("DMLC_ROLE" -> role,  
               "DMLC_PS_ROOT_URI" -> schedulerHost,  
               "DMLC_PS_ROOT_PORT" -> schedulerPort,  
               "DMLC_NUM_WORKER" -> numWorker,  
               "DMLC_NUM_SERVER" -> numServer)
```

```
KVStoreServer.init(envs)
```


Start scheduler and servers  
on different nodes

```
if (role == "server" || role == "scheduler") {  
  // scheduler & server  
  KVStoreServer.start()  
} else {  
  // worker  
  val kv = KVStore.create("dist_sync")  
  model.fit(trainData = train, kvStore = kv)  
}
```

BSP strategy. Or 'dist\_async' for fully  
asynchronous update.

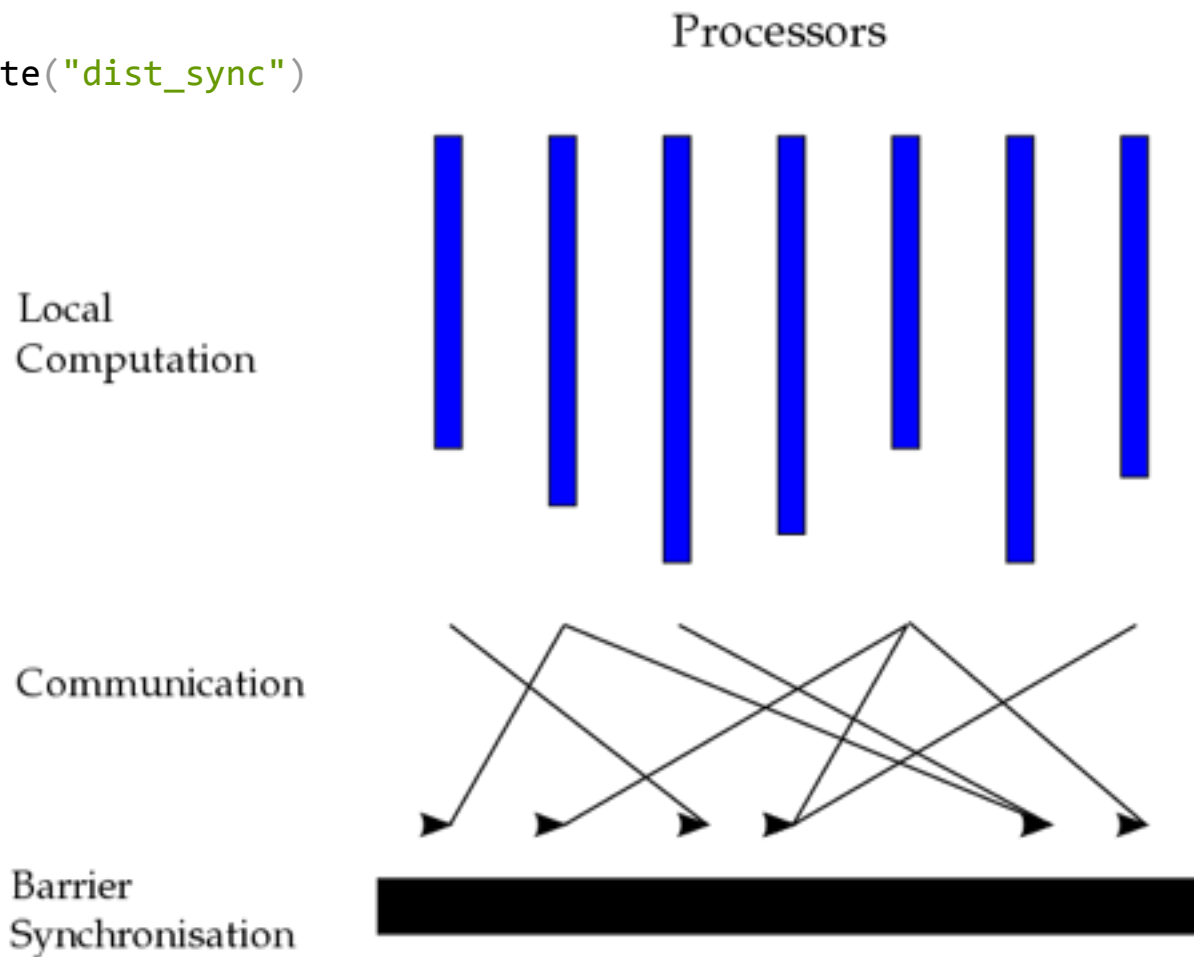
Run model fitting on  
worker node

## Parameter Server




# Bulk Synchronous Parallel

```
KVStore.create("dist_sync")
```



# Asynchronous Execution



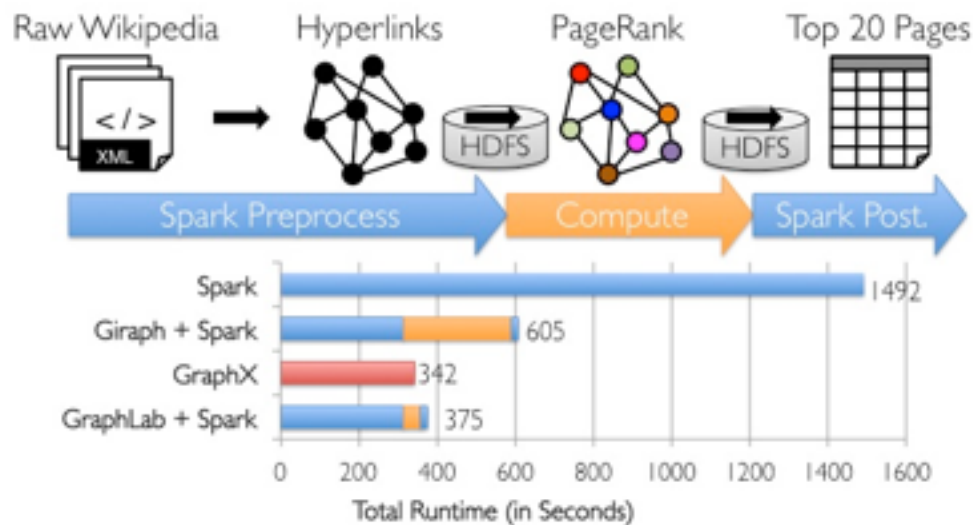
```
KVStore.create("dist_async")
```



## Why Spark

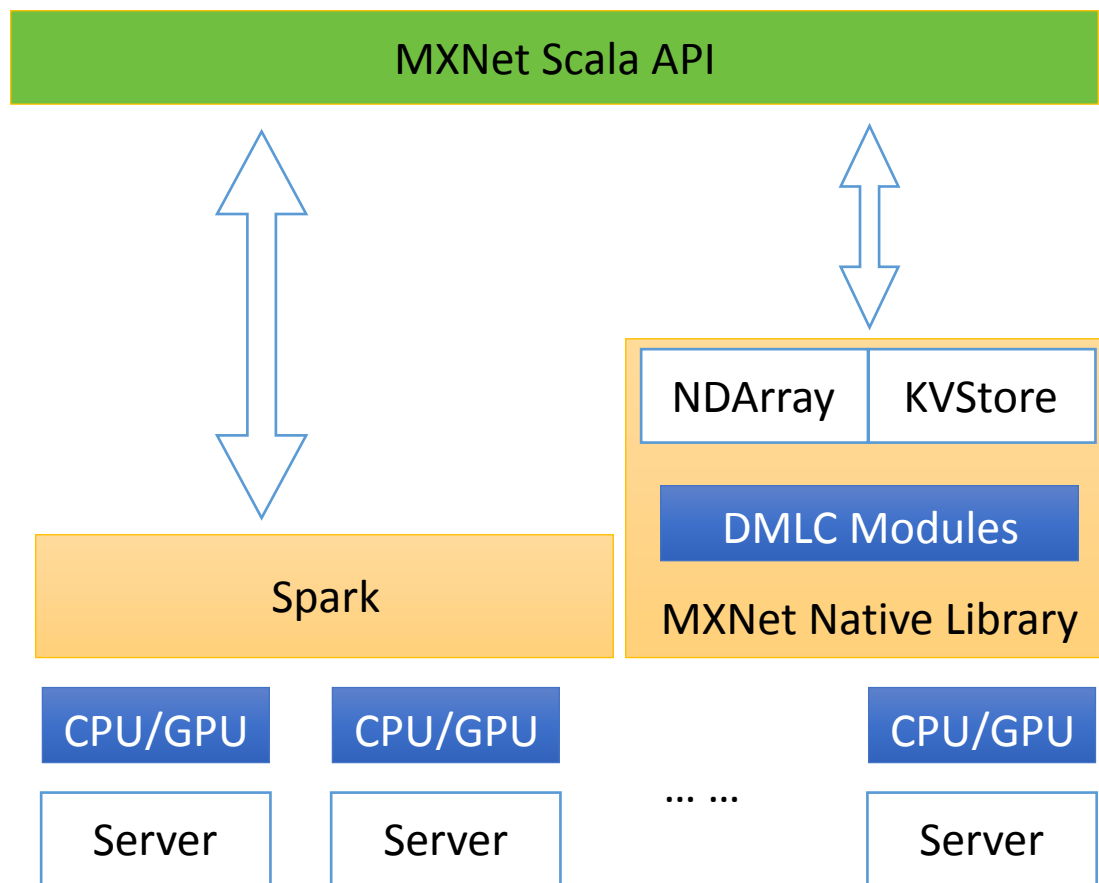
- Spark has become the de facto standard for large-scale data processing.
- Combine ETL with machine learning pipeline.

### A Small Pipeline in GraphX



Timed end-to-end GraphX is *faster* than GraphLab

## MXNet on Spark



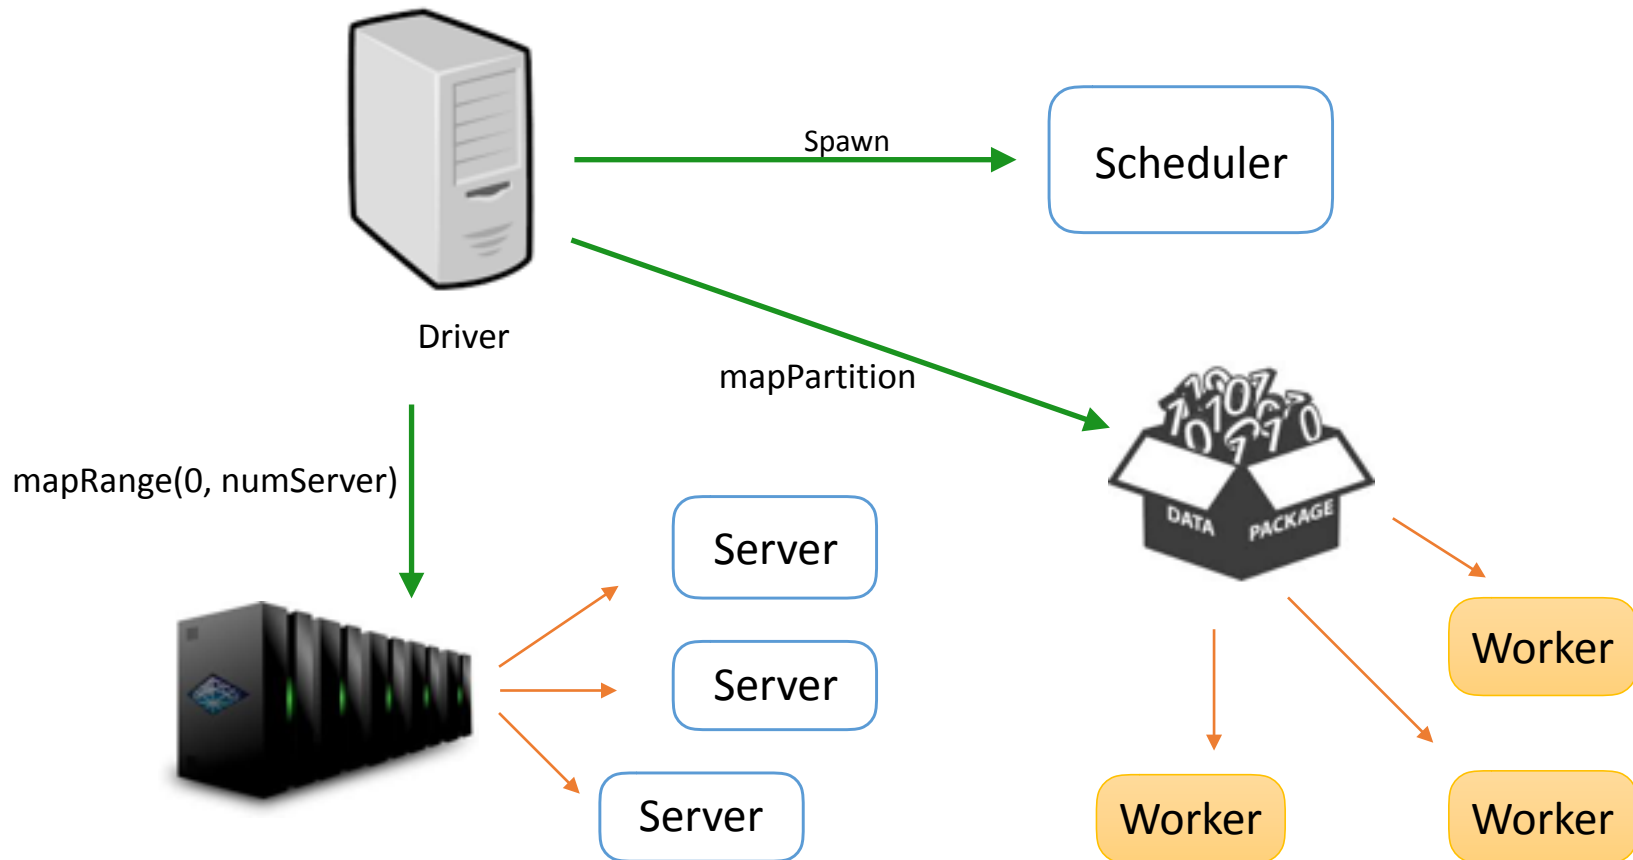
## MXNet on Spark API

```
val mxnet = new MXNet()  
  .setBatchSize(128)  
  .setLabelName("softmax_label")  
  .setContext(Context.gpu(0))  
  .setDimension(Shape(784))  
  .setNetwork(mlp)  
  .setNumServer(2)  
  .setNumWorker(4)  
  .setExecutorClasspath(classpaths)  
  .setOptimizer(  
    new SGD(learningRate = 0.01f, momentum = 0.9f, wd = 0.00001f)  
  )  
  
val model = mxnet.train(trainData)  
model.save(modelPath)  
  
val predictions = model.predict(testData)
```

mlp will be serialized to workers

setup server & worker number


## Parameter Server Components on Spark




## Performance

■ MXNet GPU   ■ MXNet CPU   ■ Spark ML

Train 660,000 samples on 5 EC2 nodes

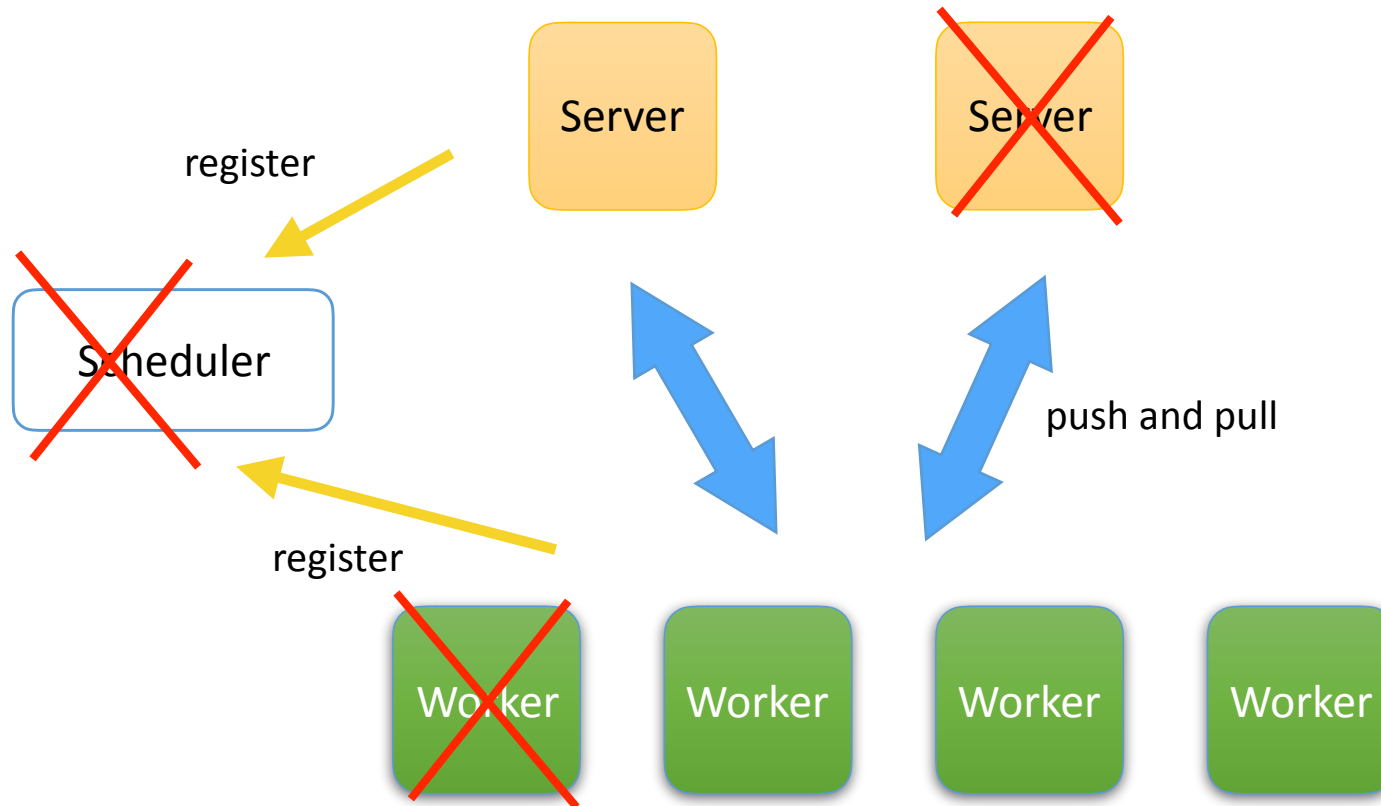


## Executors and Processes



while parameter server and mx engine is singleton

## Failover





## Resource Management

- CPU & GPU allocation on cluster
- What if there's no enough resources in cluster?

how to allocate?

GPU 0

GPU 1

GPU 2

GPU 3

```
% create executor for each GPU
execs = [symbol.bind(mx.gpu(i)) for i in range(ngpu)]
% w -= learning_rate * grad
kvstore.set_updater(...)
% iterating on data
for dbatch in train_iter:
    % iterating on GPUs
    for i in range(ngpu):
        % read a data partition
        copy_data_slice(dbatch, execs[i])
        % pull the parameters
        for key in update_keys:
            kvstore.pull(key, execs[i].weight_array[key])
        % compute the gradient
        execs[i].forward(is_train=True)
        execs[i].backward()
        % push the gradient
        for key in update_keys:
            kvstore.push(key, execs[i].grad_array[key])
```

# XGBoost

Instance index      gradient statistics


1  g1, h1

2  g2, h2

3  g3, h3

4  g4, h4


5  g5, h5




$$Obj = - \sum_j \frac{G_j^2}{H_j + \lambda} + 3\gamma$$

The smaller the score is, the better the structure is


## XGBoost Single Machine Performance



## XGBoost Distributed Training Performance

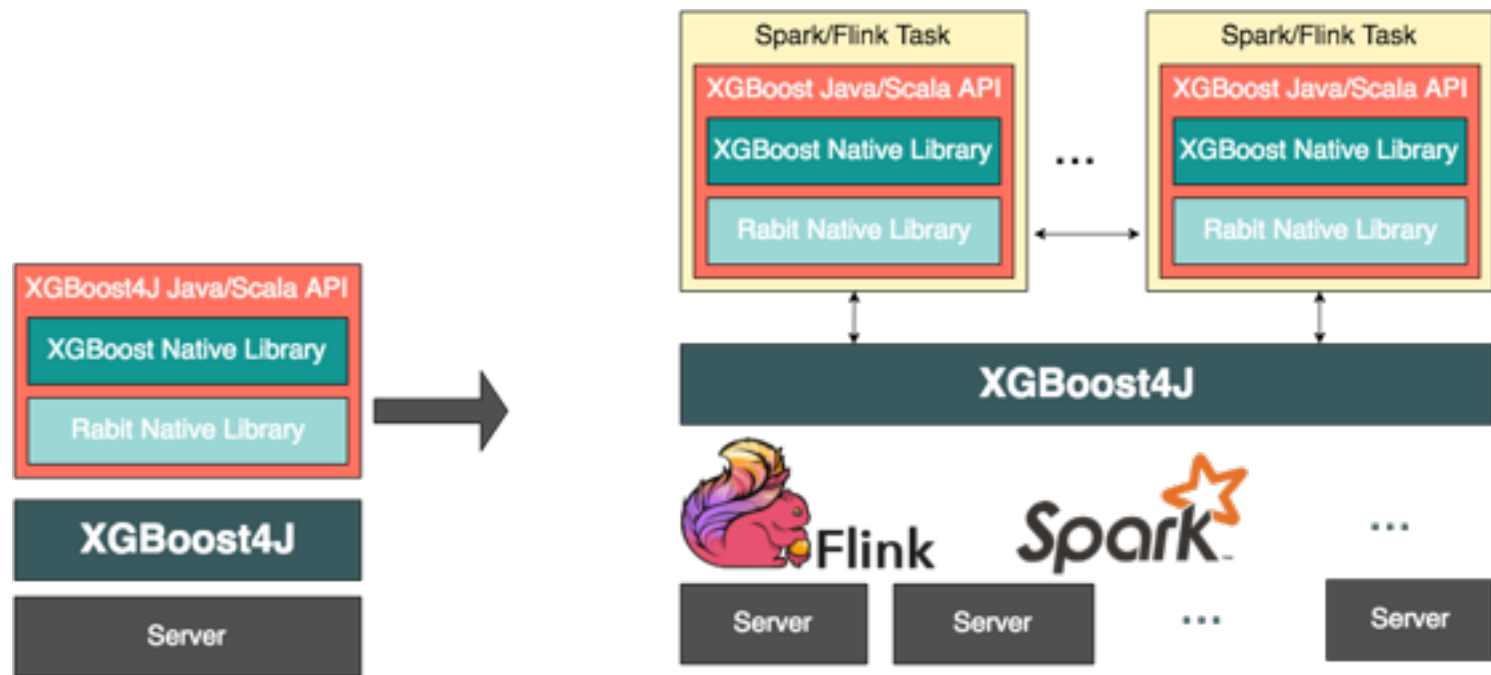


(a) End-to-end time cost include data loading



(b) Per iteration cost exclude data loading

# XGBoost on Spark



# XGBoost Training on Single Machine

```
val params = new mutable.HashMap[String, Any]()
params += "eta" -> 1.0
params += "max_depth" -> 2
params += "silent" -> 1
params += "objective" -> "binary:logistic"

val watches = new mutable.HashMap[String, DMatrix]
watches += "train" -> trainMax
watches += "test" -> testMax

val round = 2
// train a model
val booster = XGBoost.train(trainMax, params.toMap, round, watches.toMap)

val predicts = booster.predict(testMax)
```

## XGBoost Training on Spark

parse libsvm and wrap it as RDD

or, DataSet in Flink

```
val trainRDD  
    = MLUtils.loadLibSVMFile(sc, inputTrainPath).repartition(args(1).toInt)  
val xgboostModel = XGBoost.train(trainRDD, paramMap, numRound, numWorkers)
```

```
// testSet is an RDD containing testset data represented as  
// org.apache.spark.mllib.regression.LabeledPoint
```

```
val testSet = MLUtils.loadLibSVMFile(sc, inputTestPath)
```

```
// local prediction
```

```
// import methods in DataUtils to convert
```

```
Iterator[org.apache.spark.mllib.regression.LabeledPoint]
```

```
// to Iterator[ml.dmlc.xgboost4j.LabeledPoint] in automatic
```

```
import DataUtils._
```

```
xgboostModel.predict(new DMatrix(testSet.collect().iterator)
```

```
// distributed prediction
```

```
xgboostModel.predict(testSet)
```

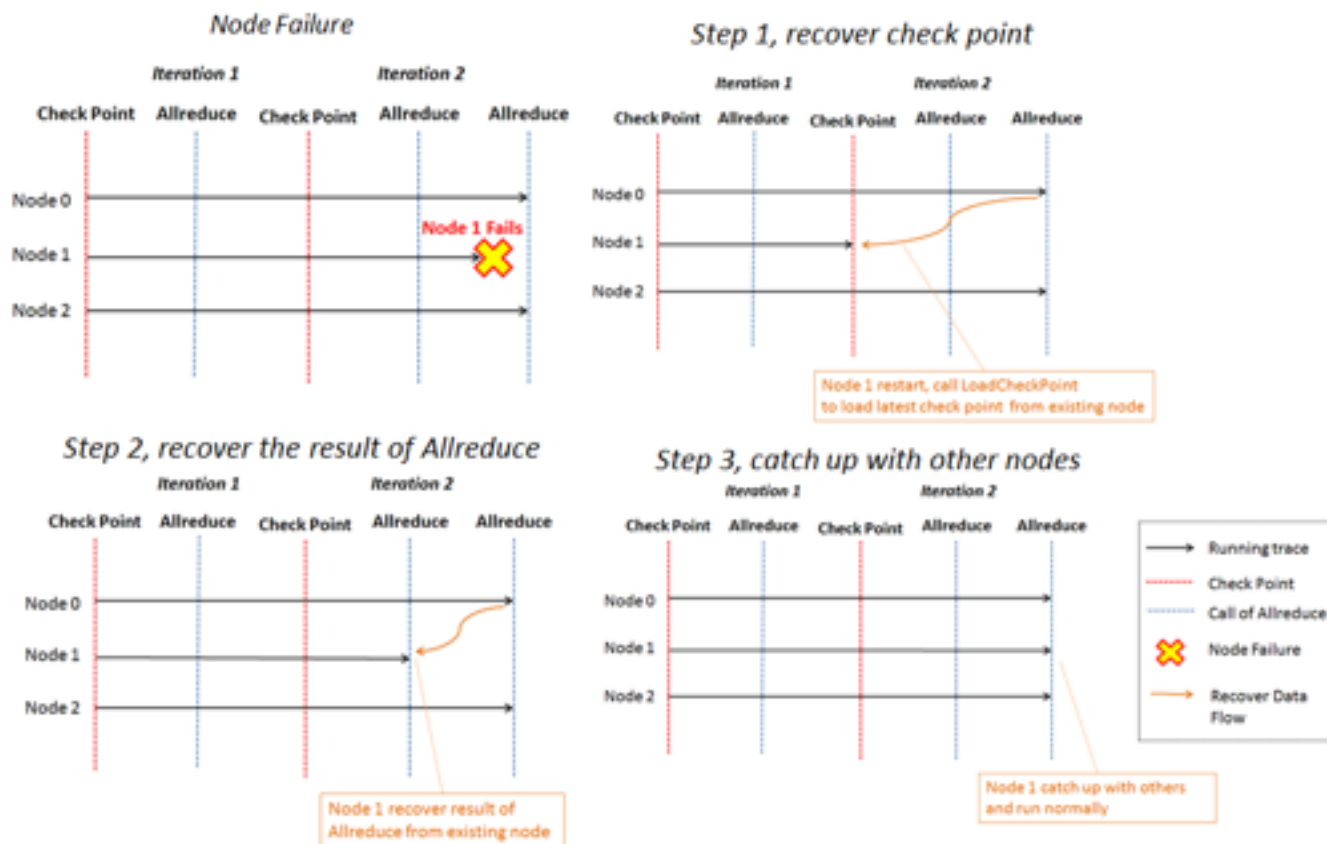
predict locally



## Rabit Allreduce

- Portable: rabbit is light weight and runs everywhere.
  - Rabbit is a library instead of a framework.
  - Rabbit only relies on a mechanism to start program, which was provided by most framework.
- Scalable and Flexible: rabbit runs fast
  - Rabbit program use Allreduce to communicate, and do not suffer the cost between iterations of MapReduce abstraction.
  - Programs can call rabbit functions in any order, as opposed to frameworks where callbacks are offered and called by the framework.
  - Programs persist over all the iterations, unless they fail and recover.
- Reliable: rabbit dig burrows to avoid disasters
  - Rabbit programs can recover the model and results using synchronous function calls.

# Rabit Failover



## Conclusion

- DMLC is committed to build efficient, flexible and portable machine learning systems.
- Our solution takes advantage of the flexible parallel training approaches and GPU support with DMLC's underlying modules, and the fast data processing pipeline with Spark. We combine the strengths of both side to scale Boosting Trees to larger dataset and faster converge rate, and bring large-scale distributed deep learning to Spark.
- Our JVM stack solution for XGBoost/MXNet is universal, it also works for other data processing systems, which means users can integrate it into their product pipeline easily.

## Roadmap

- Parameter server failure tolerance
- Run multiple tasks in one process
- Resource management on cluster
- Improve input data iterators
- Cross validation and parameter selection
- More convenient APIs
- Deploy to Maven repositories
- ... ..

## Acknowledge

- Tianqi Chen: initiator of XGBoost and MXNet.
- Mu Li: guy behind the ps-lite and kvstore.
- Nan Zhu: creator of XGBoost on Spark.
- Zixuan Huang: contributor of Java/Scala package for XGBoost and MXNet.
- Yuan Tang: contributor of Scala package for MXNet.
- Hundreds of contributors:
  - <https://github.com/dmlc/mxnet/blob/master/CONTRIBUTORS.md>
  - <https://github.com/dmlc/xgboost/blob/master/CONTRIBUTORS.md>

# Looking for contributors

<https://github.com/dmlc>

Thank you